

**Assessing Fracture Healing  
with Artificial Intelligence:**

Using Transfer Learning to Predict the  
*Radiographic Union Score for Tibial Fractures*,  
in the Radiography of High-Energy Trauma

Shen Zhou Hong      Goldsmiths, UoL

April 4th, 2023

# Contents

<b>1</b>	<b>Implementation</b>	<b>2</b>
1.1	K-Fold Evaluation . . . . .	2
1.2	Establishing a Baseline . . . . .	4
1.2.1	Shallow Convolutional Neural Network . . . . .	4
1.2.2	End-to-End Training with InceptionV3 . . . . .	6
1.2.3	Baseline Metrics . . . . .	8
1.3	InceptionV3 with Transfer Learning . . . . .	8
1.3.1	Base Model Trained on RadImageNet Dataset . . . . .	8
1.3.2	Base Model Trained on InceptionV3 Dataset . . . . .	9
1.3.3	Comparison between RadImageNet and ImageNet . . . . .	10
1.4	Hyperparameter Search . . . . .	10
1.4.1	Hyperparameter Search Regime I . . . . .	10
1.4.2	Hyperparameter Search Regime II . . . . .	14
1.4.3	Final Hyperparameters . . . . .	14
1.5	Final Model Performance . . . . .	14
<b>A</b>	<b>Additional Materials</b>	<b>15</b>
A.1	Project Code and Github Repository . . . . .	15
A.1.1	Initial Evaluation Models . . . . .	15
A.1.2	Hyperparameter Search Code . . . . .	15
A.1.3	Analysis Notebooks . . . . .	15

# Chapter 1

## Implementation

### 1.1 K-Fold Evaluation

---

```
def k_fold_dataset(ds: tf.data.Dataset, k: int = 10) -> list[tuple[tf.data.Dataset,
↳ tf.data.Dataset]]:
    # First shard the given dataset into k individual folds.
    list_of_folds: list[tf.data.Dataset] = []
    for i in range(k):
        fold: tf.data.Dataset = ds.shard(num_shards=k, index=i)
        list_of_folds.append(fold)

    # Next, generate a list of train and validation dataset tuples
    list_of_ds_pairs: list[tuple[tf.data.Dataset, tf.data.Dataset]] = []
    for i, holdout_fold in enumerate(list_of_folds):
        ds_valid: tf.data.Dataset = holdout_fold

        # Select every fold except holdout_fold as the training folds
        training_folds: list[tf.data.Dataset] = list_of_folds[:i] +
↳ list_of_folds[i+1:]

        # ds_train size is  $\frac{k-1}{k}$  of the original dataset
        ds_train: tf.data.Dataset = training_folds[0]
        for fold in training_folds[1:]:
            ds_train = ds_train.concatenate(fold)

        ds_pair: tuple[tf.data.Dataset, tf.data.Dataset] = (ds_train, ds_valid)
        list_of_ds_pairs.append(ds_pair)

    return list_of_ds_pairs
```

---

Listing 1: Sharding dataset for K-Fold Cross Validation ([Github](#))

---

```

def cross_validate(ModelClass: tf.keras.Model, ds: tf.data.Dataset, epochs: int = 50,
↳ batch_size: int = 128, k: int = 10) -> list[tf.keras.callbacks.History]:

    history_list: list[tf.keras.callbacks.History] = []
    train_valid_pairs: list[tf.data.Dataset] = k_fold_dataset(ds, k)

    for i, (ds_train, ds_valid) in enumerate(train_valid_pairs):

        tf.keras.backend.clear_session()
        model = ModelClass()
        model.compile(
            optimizer=tf.keras.optimizers.Adam(),
            loss=tf.keras.losses.BinaryCrossentropy(),
            metrics=metrics
        )
        history = model.fit(
            ds_train,
            validation_data=ds_valid,
            epochs=epochs,
            batch_size=batch_size,
        )
        history_list.append(history.history)

    return history_list

```

---

Listing 2: K-Fold Cross Validation ([Github](#))

---

```

def calculate_mean_metrics(kfold_metrics: list[dict[str, float]]) -> dict[str,
↳ list[float]]:
    # Initialise aggregate metrics with appropriate keys
    aggregate_metrics: dict[str, list[float]] = {}
    for fold in kfold_metrics:
        for metric in fold.keys():
            if metric not in aggregate_metrics:
                aggregate_metrics[metric] = []

    # Calculate the average metric per epoch for every fold
    number_of_folds: int = len(kfold_metrics)
    for metric in aggregate_metrics.keys():
        number_of_epochs: int = len(kfold_metrics[0][metric])
        for epoch in range(number_of_epochs):
            # A list of every value for that given metric in this epoch across folds
            values_per_epoch: list[float] = [x[metric][epoch] for x in kfold_metrics]
            mean_per_epoch : float = sum(values_per_epoch) / number_of_folds
            aggregate_metrics[metric].append(mean_per_epoch)

    return aggregate_metrics

```

---

Listing 3: Calculating Mean Metrics from K-Fold Data ([Github](#))

## 1.2 Establishing a Baseline

### 1.2.1 Shallow Convolutional Neural Network

---

```
class LeNet1998(tf.keras.Model):
    def __init__(self, **kwargs):
        super().__init__(**kwargs)

        self.input_layer: tf.Tensor = layers.InputLayer(input_shape=(299, 299, 3))
        self.data_augmentation: tf.keras.Sequential = tf.keras.Sequential([
            layers.RandomFlip(seed=RNG_SEED),
        ])

        self.lenet1999: tf.keras.Model = tf.keras.Sequential([
            layers.Conv2D(6, kernel_size=5, strides=1, activation='tanh',
                ⇨ padding='same'),
            layers.AveragePooling2D(),
            layers.Conv2D(16, kernel_size=5, strides=1, activation='tanh',
                ⇨ padding='valid'),
            layers.AveragePooling2D(),
        ])

        self.classifier: tf.keras.Sequential = tf.keras.Sequential([
            layers.Flatten(),
            layers.Dense(1024, activation='relu'),
            layers.Dense(18, activation='sigmoid')
        ])

        self.model: tf.keras.Sequential = tf.keras.Sequential([
            self.input_layer,
            self.data_augmentation,
            self.lenet1999,
            self.classifier
        ])

    def call(self, inputs):
        return self.model(inputs)
```

---

Listing 4: The LeNet 1998 Shallow CNN Model ([Github](#))

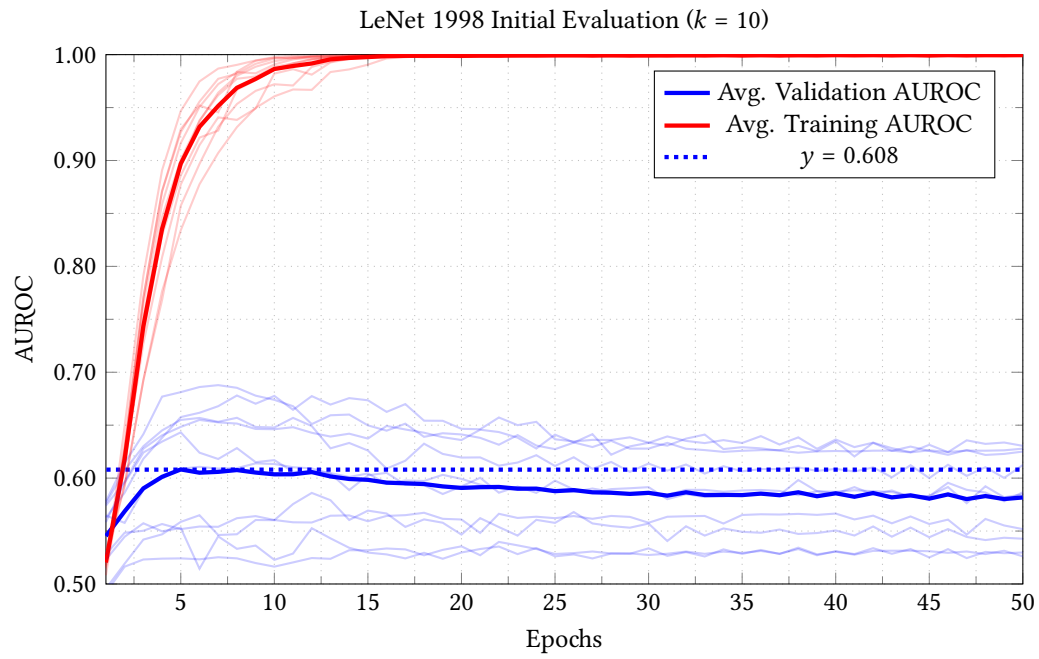


Figure 1.1: Baseline shallow CNN based on the LeNet 1998 architecture

## 1.2.2 End-to-End Training with InceptionV3

---

```
class TransferLearningModel(tf.keras.Model):
    def __init__(self, dropout_rate: float, **kwargs):
        super().__init__(**kwargs)

        self.input_layer: tf.Tensor = layers.InputLayer(input_shape=(299, 299, 3))
        self.data_augmentation: tf.keras.Sequential = tf.keras.Sequential([
            layers.RandomFlip(seed=RNG_SEED),
        ])

        self.inceptionv3: tf.keras.Model = tf.keras.applications.InceptionV3(
            include_top=False,
            weights='imagenet'
        )
        self.inceptionv3.trainable = False

        self.classifier: tf.keras.Sequential = tf.keras.Sequential([
            layers.GlobalMaxPooling2D(),
            layers.Dense(1024, activation='relu'),
            layers.Dropout(dropout_rate),
            layers.Dense( 512, activation='relu'),
            layers.Dropout(dropout_rate),
            layers.Dense( 256, activation='relu'),
            layers.Dropout(dropout_rate),
            layers.Dense( 18, activation='sigmoid')
        ])

        self.model: tf.keras.Sequential = tf.keras.Sequential([
            self.input_layer,
            self.data_augmentation,
            self.inceptionv3,
            self.classifier
        ])

    def call(self, inputs):
        return self.model(inputs)
```

---

Listing 5: Model Class for InceptionV3 ([Github](#))

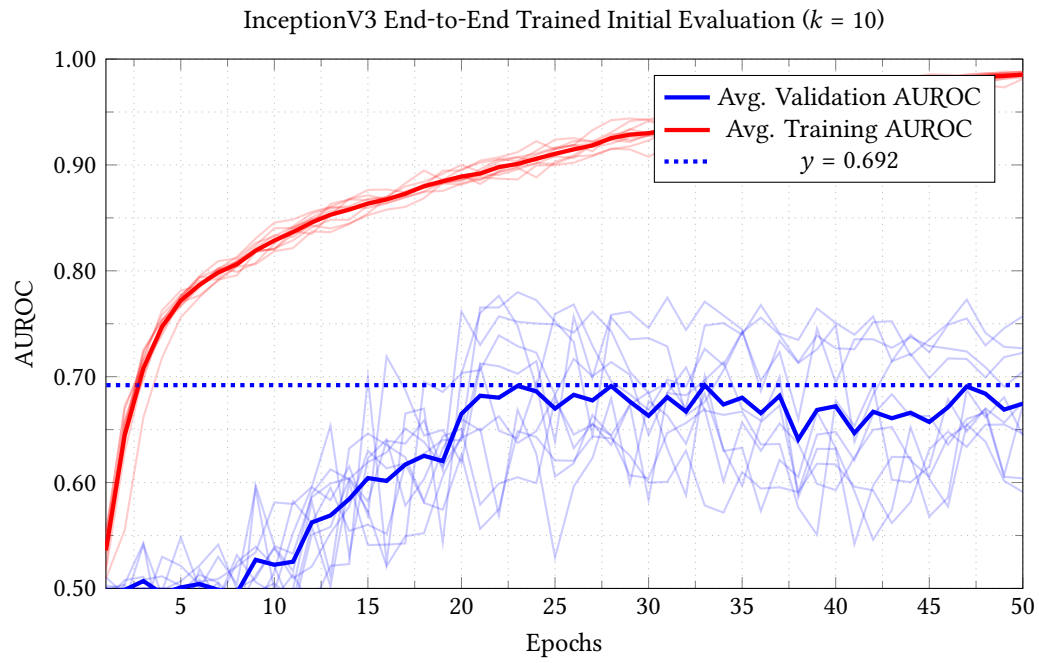


Figure 1.2: InceptionV3 Model Trained on Study Data.



### 1.2.3 Baseline Metrics

## 1.3 InceptionV3 with Transfer Learning

### 1.3.1 Base Model Trained on RadImageNet Dataset

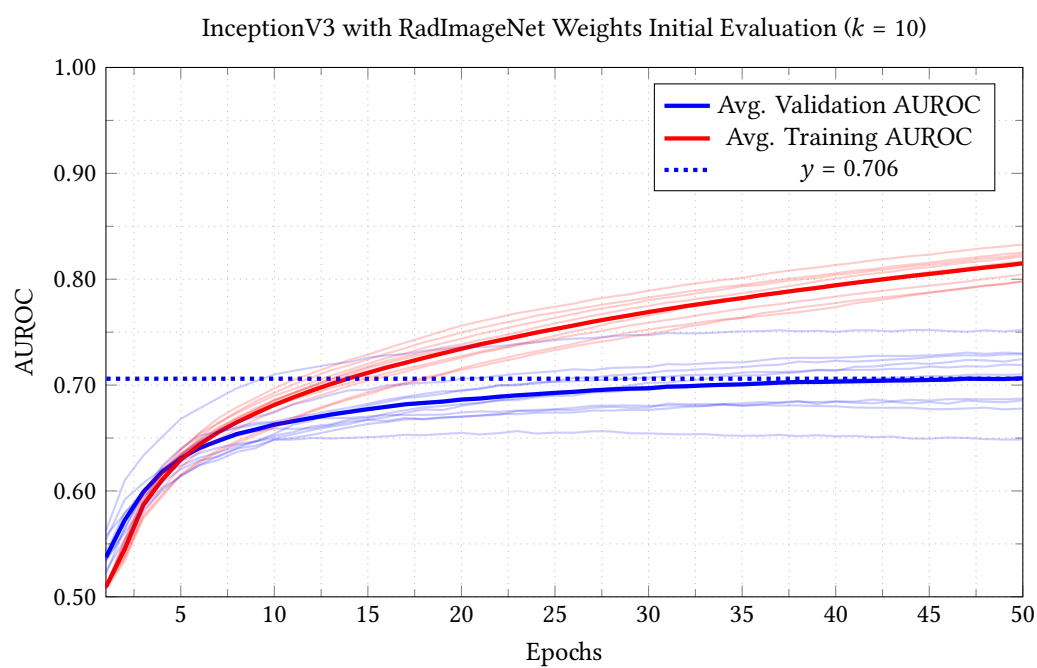


Figure 1.3: InceptionV3 with RadImageNet Weights

### 1.3.2 Base Model Trained on InceptionV3 Dataset

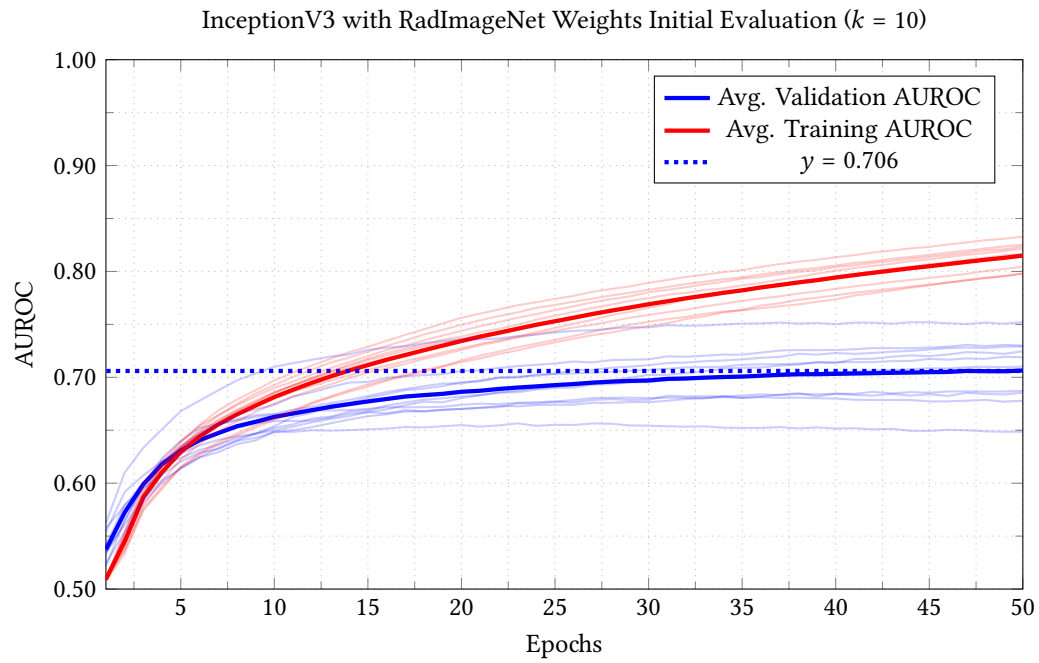


Figure 1.4: InceptionV3 with RadImageNet Weights

### 1.3.3 Comparison between RadImageNet and ImageNet

## 1.4 Hyperparameter Search

### 1.4.1 Hyperparameter Search Regime I

---

```
def hyperparameter_search(trials: int, kfold: int = 6, epochs: int = 20) ->
    list[dict[str, Union[int, float, list[tf.keras.callbacks.History]]]]:
    search_results: list[dict[str, any]] = []

    for trial in range(trials):
        # Randomly pick hyperparameter options
        rng = np.random.default_rng()
        batch_size: int = rng.integers(16, 2048, endpoint=True)
        dropout_rate: float = rng.uniform(0.0, 0.5)

        # Conduct K-Fold cross-validation with given hyperparameters
        results: list[tf.keras.callbacks.History] = cross_validate(
            TransferLearningModel,
            ds_train_and_valid,
            epochs=epochs,
            batch_size=batch_size,
            dropout_rate=dropout_rate,
            k=kfold
        )

        search_results.append({
            "max_val_auc": calc_kfold_max(results, "val_auc"),
            "batch_size": batch_size,
            "dropout_rate": dropout_rate,
            "history_list": k_fold_results
        })

    return search_results
```

---

Listing 6: Hyperparameter Search Regime I ([Github](#))

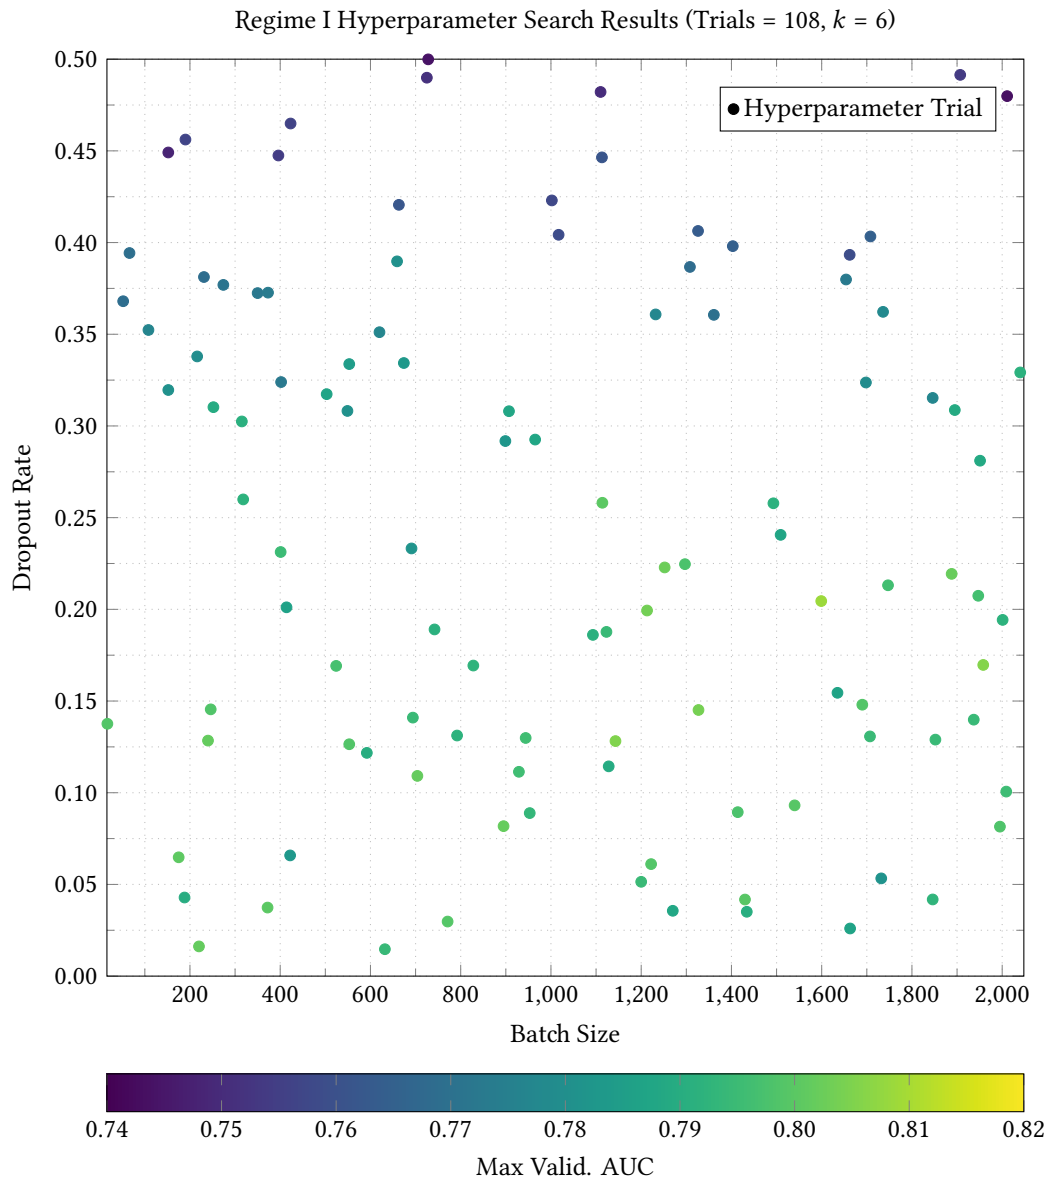


Figure 1.5: Results for the Hyperparameter Search Regime I

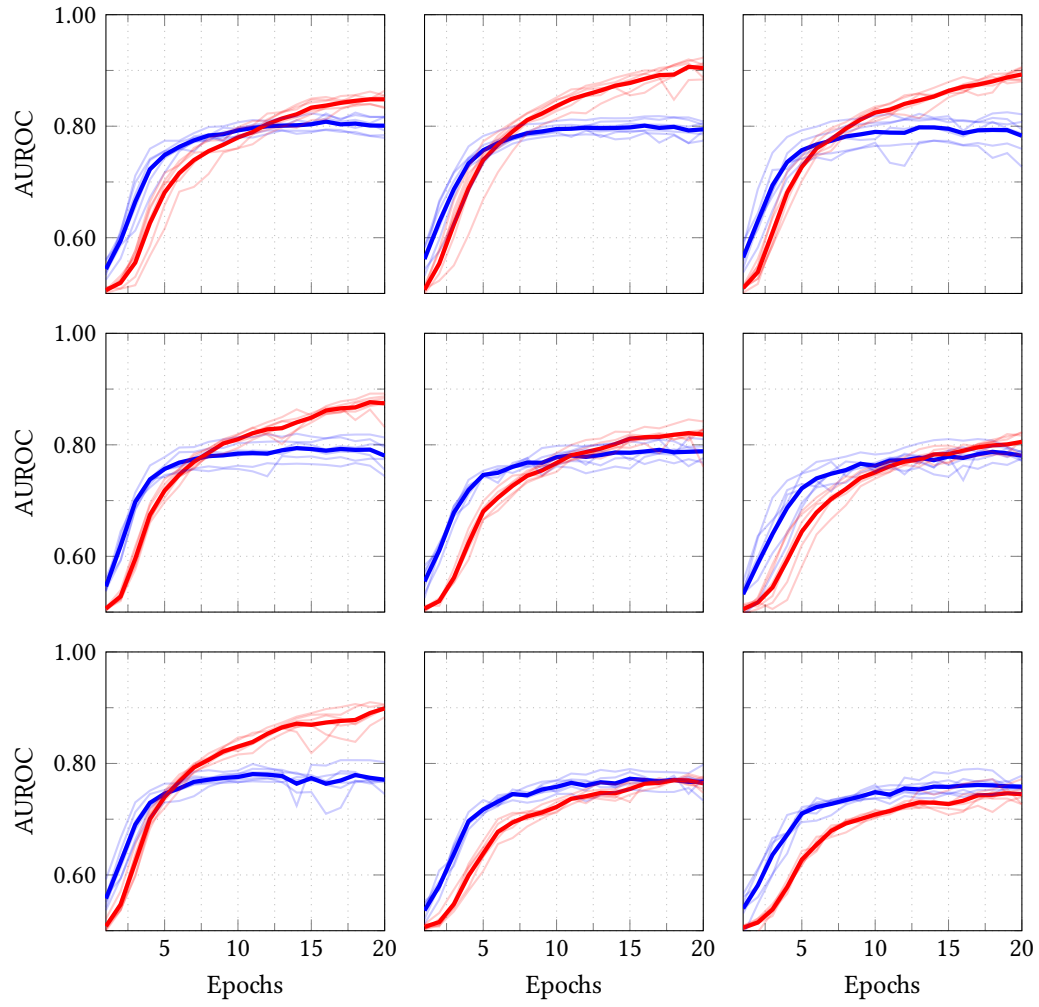


Figure 1.6: Examples of model performance from hyperparameter regime I search.

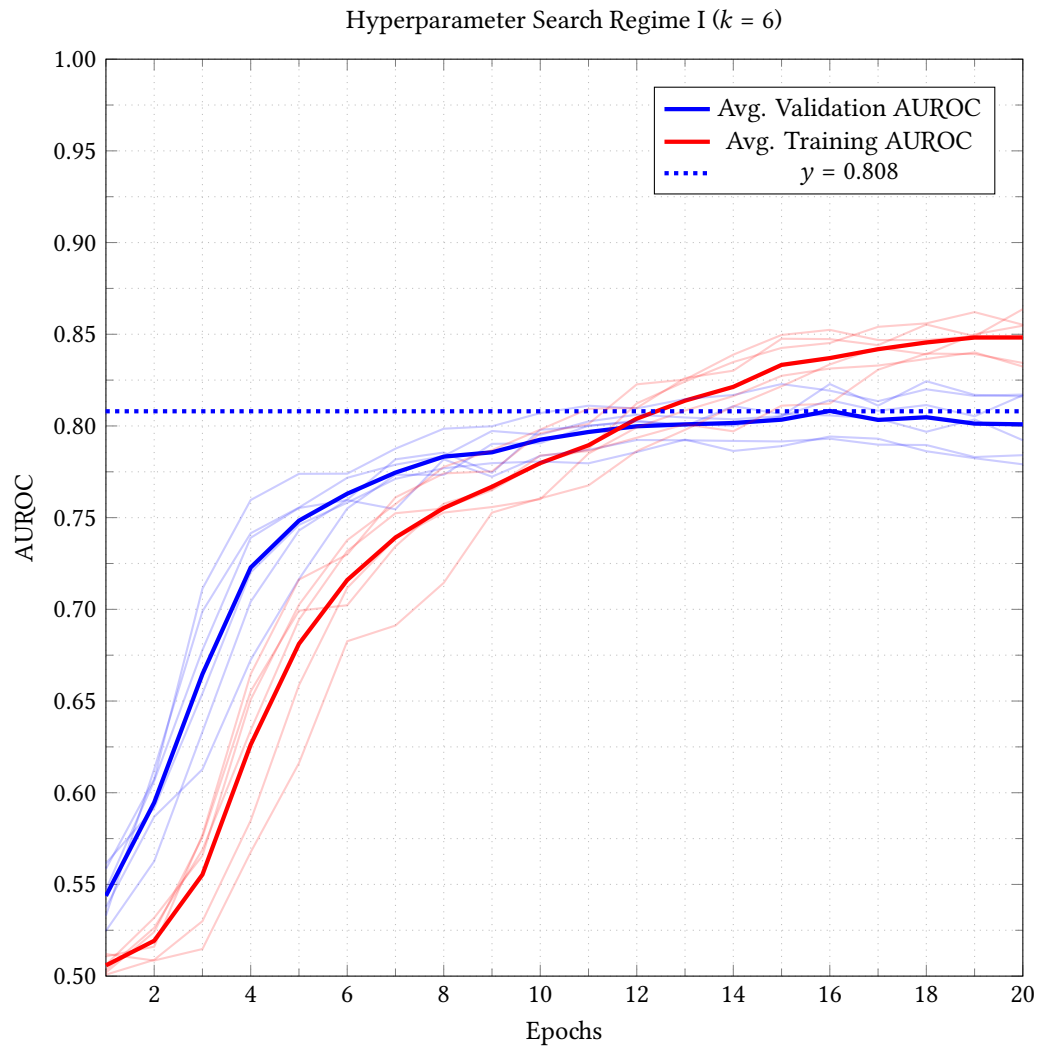


Figure 1.7: Best performing model in Regime I

### 1.4.2 Hyperparameter Search Regime II

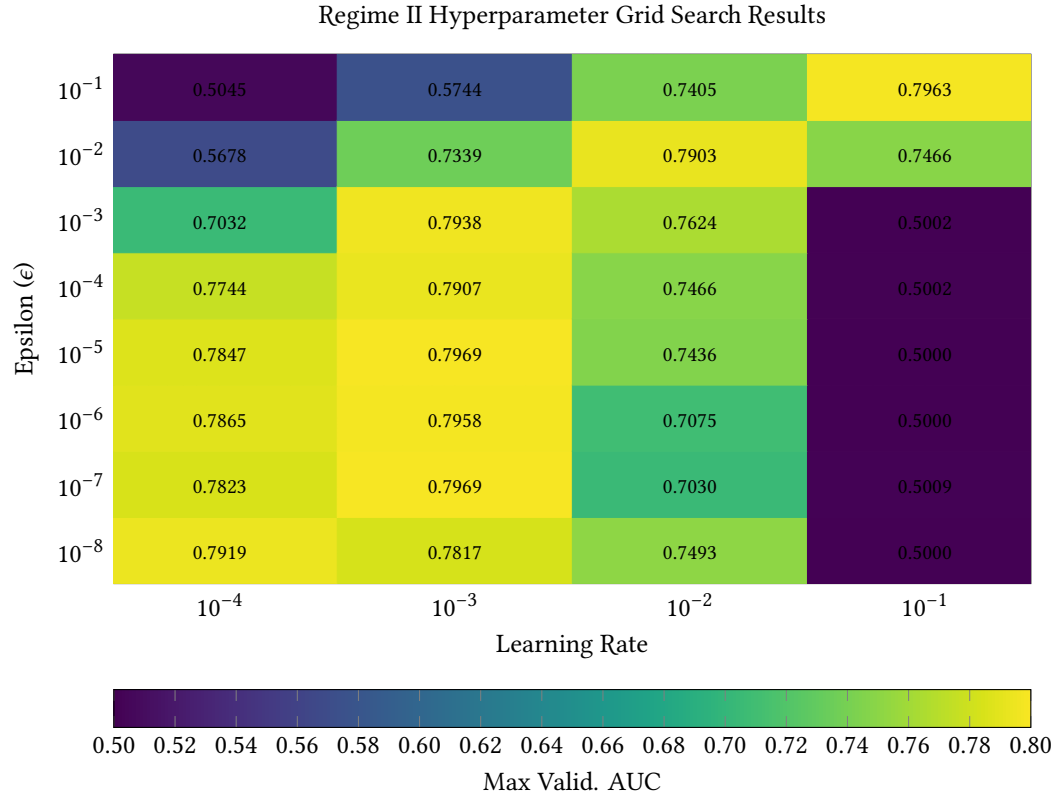


Figure 1.8: Results for the Hyperparameter Search Regime II

### 1.4.3 Final Hyperparameters

## 1.5 Final Model Performance

# Appendix A

## Additional Materials

### A.1 Project Code and Github Repository

All of the Python code used in this project (including experiment and analysis code) are available within the project Git repository, hosted on [Github](#). The code is located within the `python/` directory of the repository root:

<https://github.com/ShenZhouHong/radiography-ai-project/>

#### A.1.1 Initial Evaluation Models

Jupyter notebooks used to run the initial evaluations of LeNet 1998, InceptionV3 with end-to-end training, and initial transfer learning models:

<https://github.com/ShenZhouHong/radiography-ai-project/tree/master/python/initial-evaluation>

#### A.1.2 Hyperparameter Search Code

Jupyter notebooks used to perform the hyperparameter search regime.

<https://github.com/ShenZhouHong/radiography-ai-project/tree/master/python/hyperparam-search>

#### A.1.3 Analysis Notebooks

Jupyter notebooks used to analyse the raw data, process for insights and visualisations, and output CSV files:

<https://github.com/ShenZhouHong/radiography-ai-project/tree/master/python/analysis>